




Systems Design: Multi-Cloud MLOps Pipeline



"Deploy ML models across AWS, GCP, and Azure
without vendor lock-in"

Where do you start?

Most people pick one cloud and pray it works
I design for portability from day one.

Here's the thinking process: 



Step 1: Why Multi-Cloud?



Why not just
pick AWS and be
done?

Because:

- **Vendor lock-in** = pricing power (**they own you**)
- **Single cloud** = **single point of failure**
- **Best-of-breed** = different clouds excel at different things

The real question:


"How do we deploy models anywhere **WITHOUT** rewriting everything?"

This question **determines your architecture.**

Wrong approach = locked in forever.

Right approach = deploy anywhere in hours.

⟳ Step 2: Map the MLOps Flow



Training → Registry → Deploy → Monitor

Fill in the [?]:

- Where do models train? (Which cloud? Both?)
- Where do artifacts live? (One registry? Three?)
- How do we deploy? (Cloud-native? Containers?)
- How do we monitor? (One dashboard? Three?)

This is the **skeleton**.

Now add the **organs** that make it **cloud-agnostic**.

Step 3: Core Components

Training (**Kubernetes**, **Kubeflow**, or **cloud-native**)



Model Registry (**MLflow**, **W&B**, **cloud-native**)



Deployment (**SageMaker**, **Vertex**, **Azure ML**, or
K8s)



Monitoring (**Prometheus**, **CloudWatch**, **unified dashboard**)

But this is **TOO** simple.

What's missing?

- How do we **move models** between clouds?
- How do we **handle cloud-specific** APIs?
- What if a **cloud goes down**?
- How do we **track costs** across three clouds?

Add the **layers** that multi-cloud requires.



Step 4: Multi-Cloud Artifact Registry

Cloud-Agnostic Model Registry:

Option 1: MLflow (self-hosted)

- Works anywhere
- Cloud-agnostic
- You manage infrastructure

Option 2: Weights & Biases

- SaaS (no infrastructure)
- Multi-cloud native
- Unified tracking

Option 3: Cloud-native + sync

- AWS S3 ↔ GCP GCS ↔ Azure Blob
- Sync artifacts across clouds
- More complex, full control

Questions to answer:

- **Single source of truth?** (Yes—one registry)
- **Versioning strategy?** (Semantic versioning)
- **Metadata?** (Training config, metrics, lineage)

This is where your **models live**.

Get this **wrong** = can't **deploy** anywhere.

Cloud-Specific Deployment Layers

AWS Deployment:

- SageMaker (managed inference)
- EKS (Kubernetes)
- Lambda (serverless)

GCP Deployment:

- Vertex AI (managed)
- GKE (Kubernetes)
- Cloud Run (serverless)

Azure Deployment:

- Azure ML (managed)
- AKS (Kubernetes)
- Azure Functions (serverless)

The trick: Abstraction layer

Container-based deployment (Docker):

- Build once
- Deploy to SageMaker, Vertex, Azure ML
- Or deploy to K8s (EKS, GKE, AKS)

Kubernetes = your portability layer.

Cloud APIs change. Containers don't.

Step 6: Unified Monitoring Across Clouds

The Problem:

- AWS CloudWatch (only sees AWS)
- GCP Cloud Monitoring (only sees GCP)
- Azure Monitor (only sees Azure)

You need: **ONE dashboard, THREE clouds**

Solution: Cloud-agnostic monitoring

Prometheus + Grafana:

- Scrapes metrics from all clouds
- Unified dashboards
- Consistent alerting

Track:

- Inference latency (per cloud)
- Model drift (unified across clouds)
- Error rates (compare clouds)
- Costs (which cloud is expensive?)

Without unified monitoring = flying blind across three clouds.

You can't optimize what you can't see.

Step 7: CI/CD for Multi-Cloud

GitHub Actions / GitLab CI:

Push code → Run tests (unit, integration)



Build container



Push to registry (MLflow/W&B/ECR/GCR/ACR)



Deploy to: [AWS | GCP | Azure | All Three]

Key: Infrastructure as Code (Terraform)

One config file = deploy to any cloud:

→ terraform apply -var="cloud=aws"

→ terraform apply -var="cloud=gcp"

→ terraform apply -var="cloud=azure"

Secrets management:

→ HashiCorp Vault (cloud-agnostic)

→ Or cloud-native (AWS Secrets Manager, GCP Secret Manager)

CI/CD = your deployment automation.

Get this right = deploy anywhere in minutes.

⚠ Step 8: Design for Cloud-Specific Failures

Where does multi-cloud break?

Cloud-specific quirks:

- AWS: IAM permissions (complex, service-specific)
- GCP: Service accounts (different auth model)
- Azure: Resource groups (deployment scope)

Network failures:

- Cross-cloud latency (AWS → GCP = 50-100ms)
- Egress costs (moving data OUT = expensive)

Cost surprises:

- AWS: Data transfer out (\$0.09/GB)
- GCP: Networking egress (\$0.08/GB)
- Azure: Bandwidth (\$0.087/GB)

Then design defenses:

- Cloud-agnostic abstractions (Terraform, K8s)
- Fallback deployments (if AWS fails, route to GCP)
- Cost monitoring (alert on >\$X per cloud)
- Test BEFORE production (staging on each cloud)

Plan for cloud **failures**.

\$ Step 9: Multi-Cloud Cost Optimization

Spot instances / Preemptible VMs:

- AWS Spot: 70-90% discount
- GCP Preemptible: 80% discount
- Azure Spot: 90% discount

Why multi-cloud saves money:

Training workloads:

- Train on cheapest cloud at that moment
- Auto-switch if prices change

Inference workloads:

- Route traffic to lowest-cost region
- Or lowest latency (user experience > cost)

Storage tiering:

- Hot data: Where it's accessed most
- Cold data: Cheapest cloud (GCS Nearline, S3 Glacier)

Tools: CloudHealth, Kubecost

Track spend per cloud, per model, per team.

Multi-cloud without cost tracking = budget nightmare

The Complete Multi-Cloud MLOps System

Training (K8s or cloud-native)



Model Registry (MLflow / W&B - cloud-agnostic)



CI/CD (GitHub Actions / Terraform)



Deploy to: AWS | GCP | Azure



Unified Monitoring (Prometheus + Grafana)



Cost Tracking (per cloud, per model)

This is production **multi-cloud MLOps**.

Not "pick **CLOUD** and hope."

6 layers. Each critical.

The difference: Systems thinking.

Most teams:

- Pick one cloud
- Use cloud-native tools
- Get locked in
- Pay whatever the cloud charges

1% teams:

- Design for portability from day one
- Use cloud-agnostic tools (K8s, Terraform, MLflow)
- Deploy anywhere in hours
- Negotiate from strength (can switch clouds)

Vendor lock-in isn't technical.

It's architectural.

**Design for portability, or
pay the lock-in tax forever.**

Next in series: Petabyte-Scale Data Pipeline 