



"Design a control system that operates 550km above Earth with 240ms latency"

Where do you start?

Most people assume you can remote-control satellites like drones.

I start with the constraint: Earth is too far away.

Here's the thinking process: 



Step 1: Why Orbital is DIFFERENT

The Challenge:

Ground control → Satellite

- 240ms one-way latency (480ms round-trip)
- Connection windows: 5–15 minutes per orbit
- 16 orbits per day
- Physical inaccessibility (can't send a technician)

Questions that determine everything:

- Can you remote-control with 480ms lag? (No)
- Can you upload commands during brief windows? (Limited)
- What if connection drops mid-operation? (System must continue)
- What if something breaks? (No physical access for repairs)

Ground systems: Low latency, always-connected, accessible

Orbital systems: High latency, intermittent connection, inaccessible

Wrong approach = design like ground system, fail in orbit. Right approach = design for autonomy from day one.

Step 2: Map the Control Loop

Traditional Ground System:

Sensor → Ground Control → Decision → Command → Actuator

(Works with <50ms latency, always connected)

Orbital System:

Sensor → [240ms] → Ground → [240ms] → Actuator

(480ms round-trip = too slow for dynamic control)

Fill in the [?]:

- What decisions can ground make? (High-level only)
- What must be autonomous? (Everything time-critical)
- How do we handle disconnection? (System continues safely)
- How do we update autonomy? (Upload new logic during windows)

This is the skeleton. Now add the layers that make orbital autonomy work.



Step 3: Autonomous Control Layer

What Must Be Autonomous:

Attitude control:

- Maintain orientation (solar panels face sun)
- React to disturbances (drag, solar pressure)
- Cannot wait 480ms for ground decision

Collision avoidance:

- Detect debris (radar, visual)
- Execute avoidance maneuver (seconds to react)
- Report to ground after

Fault response:

- Sensor failure (switch to backup)
- Power loss (enter safe mode)
- Thermal anomaly (adjust radiators)

Onboard Decision Logic:

Rule-based systems:

- IF battery < 20% THEN enter power-save mode
- IF debris detected < 1km THEN execute avoidance
- IF ground contact lost > 24h THEN enter safe mode

Machine learning (limited):

- Anomaly detection (is this behavior normal?)
- Predictive maintenance (will this component fail?)
- Resource optimization (power, compute, bandwidth)

The system must ASSUME it's alone. Ground control = advisor, not operator



Step 4: Intermittent Connectivity Strategy

Connection Reality:

Orbit period: 90 minutes

Ground station visibility: 5-15 minutes per pass

Passes per day: 16 (not all over friendly ground stations)

Actual contact time: ~2 hours/day (8% of time)

During Contact Window:

Downlink (Satellite → Ground):

- Telemetry (health, status, errors)
- Science data (observations, measurements)
- Logs (what happened since last contact)

Uplink (Ground → Satellite):

- High-level commands (mission objectives)
- Software updates (new autonomy logic)
- Parameter adjustments (thresholds, priorities)

Outside Contact Window (22 hours/day):

Satellite operates FULLY autonomously:

- Execute pre-planned operations
- Respond to anomalies
- Log everything for next downlink
- Prioritize what to send (bandwidth limited)

Data prioritization:

- Critical errors: Send immediately when connected
- Routine telemetry: Batch and compress
- Science data: Store, send when bandwidth available

Communication = scheduled, not real-time. Design for 92% disconnection.

⚠ Step 5: Design for Inevitable Failure Part 1

Failure Reality in Orbit:

Single event upsets (cosmic radiation flips bits):

- Happens constantly
- Must detect and correct
- ECC memory, redundant systems

Safe Mode Design:

Trigger conditions:

- Ground contact lost > 24 hours
- Battery < critical threshold
- Multiple sensor failures
- Unexpected behavior detected

Component degradation:

- Solar panels degrade (3-5% per year)
- Batteries cycle (limited charge cycles)
- Electronics age (radiation damage)

Safe mode actions:

- Point solar panels at sun (maximize power)
- Stop all non-essential operations
- Extend communication antenna (maximize contact chance)
- Wait for ground intervention

Unrepairable failures:

- Can't send a technician
- Must work around failures
- Graceful degradation

Watchdog timers:

- System must "heartbeat" every N seconds
- If heartbeat stops → automatic reboot
- If reboot fails → safe mode

Redundancy:

Triple modular redundancy (TMR):

- 3 computers vote on decisions
- Majority wins (fault tolerance)
- Can survive 1 computer failure

Failure cascade prevention:

- Isolate faults (don't let one failure kill everything)
- Graceful degradation (lose capability, not entire system)
- Fail-safe defaults (when uncertain, be safe)

1% orbital systems plan for failure. Because failure is guaranteed, eventually



Step 6: Resource Constraints in Orbit

Power Reality:

Daylight (45 min):

- Solar panels generate power
- Charge batteries
- Run all systems

Eclipse (45 min):

- No solar power
- Battery only
- Reduce consumption

Power budget:

- Solar panels: 200W (degrades over time)
- Battery capacity: 100Wh
- System draw: 80-150W (depending on operations)

Autonomous power management:

Low power operations (eclipse safe):

- Attitude control (minimal)
- Housekeeping
- Data storage

High power operations (daylight only):

- Data processing
- High-bandwidth transmission
- Science instruments

Emergency power (critical only):

- Survival heaters
- Core computer
- Communication receiver (listen for ground)

Thermal Reality:

Sun-facing side: +120°C

Shadow side: -150°C

Temperature swing: 270°C per orbit

Thermal management:

- Radiators (passive cooling)
- Heaters (active warming)
- Thermal mass (batteries, structure)
- Orientation control (point hot side away from sun)

Systems must survive extreme thermal cycling. And manage it autonomously.

Step 7: Update Autonomy Without Bricking

The Problem:

Software bug = satellite is now space junk

- No physical access
- No rollback if update fails
- One shot to get it right

Staged updates:

- Upload new software (may take multiple passes)
- Verify checksums (ensure no corruption)
- Ground command: "Run tests on partition B"
- Satellite reports test results
- Ground command: "Switch to partition B"
- If anything fails: Stay on partition A

Automatic rollback:

- New software boots
- Must send "I'm alive" within 10 minutes
- If no heartbeat: automatic reboot to partition A
- Ground notified of failure

Over-the-air updates:

- Differential updates (only changed bytes)
- Compressed (bandwidth limited)
- Encrypted (security)
- Signed (authenticity)

Update Strategy:

Dual-boot system:

- Partition A: Current working software
- Partition B: New software upload
- Test partition B extensively
- Only switch if tests pass

Golden image:

- Partition C: Factory default (last resort)
- Minimal functionality, guaranteed to work
- Used if both A and B corrupted

ML model updates (if applicable):

- Upload new model weights
- Test against validation set onboard
- Compare performance vs current model
- Only deploy if better

Updates are the highest-risk operation. Design for them from day one.



Step 8: Debug a System 550km Away

The Debugging Problem:

Can't attach debugger

Can't see logs in real-time

Can't reproduce environment on ground

Log prioritization:

→ Errors: Immediate downlink

→ Warnings: Next contact window

→ Info: Batch and compress

→ Debug: Only if requested

Telemetry Strategy:

What to log:

- Every state transition (mode changes)
- Every decision made (and why)
- Every anomaly detected
- Resource usage (power, memory, CPU)
- Environmental conditions (temperature, radiation)

Bandwidth-limited logging:

- Can't send everything
- Summarize (statistical aggregates)
- Sample (1% of routine data)
- Compress (lossless for critical, lossy for science)

Black box recorder:

→ Last 24 hours of data (high resolution)

→ Survives reboots, failures

→ Downloaded first if anomaly detected

Ground replay:

- High-fidelity simulator on ground
- Replay satellite telemetry
- Reproduce issues
- Test fixes before uploading

Predictive monitoring:

→ Trend analysis (battery degrading faster than expected?)

→ Anomaly detection (behavior changed?)

→ Predictive alerts (failure likely in 30 days)

You can't debug what you can't see. Log everything. Prioritize transmission.

Bonus: Why Launch East?

Earth's rotation: 1,670 km/h at equator (460 m/s eastward)

Launching EAST:

- Rocket velocity + Earth rotation = COMBINED
- Free 460 m/s boost (saves ~15% fuel)
- More payload capacity OR higher orbit

Launching WEST:

- Must overcome Earth's rotation
- Lose 460 m/s (need 920 m/s extra)
- Massive fuel penalty

This is why:

- Cape Canaveral (Florida) launches east over Atlantic
- SpaceX Starbase (Texas) launches east over Gulf of Mexico
- Baikonur (Kazakhstan) launches east
- Kourou (French Guiana) launches east over ocean
- No major launch sites face west

SpaceX chose Boca Chica, Texas, specifically for:

- Southern latitude (closer to equator = more rotational boost)
- Eastward ocean access (Gulf of Mexico)
- Minimal air traffic (south Texas remote)
- Starship orbital test flights

Orbital mechanics = constraint from second zero. You can't fight physics.

The Complete Orbital Autonomous System

Layer 1: Autonomous Decision Making
(Rule-based + ML, real-time response)

Layer 2: Intermittent Communication
(5-15 min windows, 16x/day)

Layer 3: Safe Mode & Failure Handling
(TMR, watchdogs, graceful degradation)

Layer 4: Power & Thermal Management
(45-min day/night cycles, autonomous optimization)

Layer 5: Software Updates
(Dual-boot, staged rollout, automatic rollback)

Layer 6: Observability & Debugging
(Prioritized telemetry, black box, ground replay)

Layer 7: Mission Planning & Optimization
(Multi-orbit schedules, constraint satisfaction)

Bonus: Why we launch from the East

This is orbital autonomous control. Not "remote control from ground."
7 layers of autonomy. Because Earth is 550km away.