



Systems Design: Petabyte-Scale Data Pipeline



"Ingest **10TB/day**, process it, store it, analyze it—
without breaking the bank"

Where do you start?

Most people start with "big data tools."

I start with data flow and cost constraints.

Here's the thinking process: 



Step 1: What Problem Are We Solving?

The Challenge:

10TB/day incoming data

→ **3.65 PB/year**

→ Multiple sources (APIs, databases, streams, files)

→ Different formats (JSON, CSV, Parquet, logs)

→ Need to: clean, transform, analyze, store

Questions that determine everything:

→ Real-time or batch? (Latency requirements?)

→ How long to retain? (Cost implications?)

→ Query patterns? (Analytics vs operational?)

→ Compliance? (GDPR, data residency?)

Wrong architecture = millions wasted on storage you don't need.

Right architecture = only pay for what you use.

Step 2: Map the Data Flow

Ingestion → Storage → Processing → Analytics

Fill in the [?]:

- Where does **data** come from? (APIs, DBs, streams?)
- Where does **raw data** land? (Data lake?)
- How do we **transform** it? (Spark, Airflow?)
- Where does **clean data** go? (Data warehouse?)
- How do we **query** it? (SQL, BI tools?)

This is the **skeleton**.

Now add the **medallion architecture** that production needs.

Y Step 3: Bronze → Silver → Gold

Bronze Layer (Raw):

- Data lands exactly as received
- No transformation, no validation
- Cheap storage (S3, GCS, Azure Blob)
- Append-only, never delete

Silver Layer (Cleaned):

- Validated, deduplicated
- Schema enforcement
- Type casting, null handling
- Partition by date/source

Gold Layer (Analytics-Ready):

- Aggregated, joined
- Business logic applied
- Optimized for queries
- Served to BI tools, ML models

Why three layers?

- Bronze = audit trail (what arrived?)
- Silver = reliable source (cleaned)
- Gold = fast queries (optimized)

Most teams skip Bronze. Then they can't debug.



Step 4: Data Ingestion

Sources:

- **Streaming** (Kafka, Kinesis, Pub/Sub)
- **APIs** (REST, webhooks)
- **Databases** (CDC - Change Data Capture)
- **Files** (S3, SFTP, object storage)
- **Logs** (Fluentd, Logstash)

Ingestion Tools:

- **Airbyte** (pre-built connectors)
- **Fivetran** (managed, expensive)
- **Custom** scripts (Python, Go)
- **Kafka Connect** (streaming)

Questions to answer:

- Batch or streaming? (Latency vs cost)
- Schema on read or write? (Flexibility vs validation)
- Retry logic? (What if source fails?)
- Rate limiting? (Don't overwhelm APIs)

Bronze layer = everything lands here first.
Raw, unprocessed, exactly as received.

Step 5: Processing at Scale

Processing Engine:

- Databricks (managed Spark, Delta Lake native)
- Apache Spark (open-source)
- Google Dataflow (serverless)
- AWS Glue (managed ETL)

Orchestration:

- Databricks Workflows (native scheduling)
- Airflow (DAG-based scheduling)
- Prefect (modern Airflow alternative)

Bronze → Silver transformation:

- Schema validation (reject bad data)
- Deduplication (same event twice?)
- Type casting (strings → dates → integers)
- Partitioning (by date, by source)
- Delta Lake (ACID transactions, versioning)

Silver → Gold transformation:

- Joins (combine datasets)
- Aggregations (hourly, daily rollups)
- Business logic (metrics, KPIs)
- Denormalization (optimize for reads)

Processing = where you **spend compute \$.**

Optimize here or burn budget.

Step 6: Storage Tiering

Bronze (Raw - Cheap):

- S3 Standard (frequently accessed)
- After 30 days → S3 Infrequent Access
- After 90 days → S3 Glacier (archive)
- Cost: \$0.023/GB → \$0.01/GB → \$0.004/GB

Silver (Cleaned - Medium):

- S3 Standard (queried regularly)
- Delta Lake format (ACID transactions, time travel)
- Partitioned (date, source)

Gold (Analytics - Fast):

- Databricks SQL Warehouse (interactive queries)
- Or S3 + Athena (serverless queries)
- Delta Lake optimized for BI tools

Why tier storage?

- Bronze: 3.65 PB/year @ **\$0.023/GB = \$84K/year**
- Bronze with tiering: ~**\$40K/year (50% savings)**

Lifecycle policies = automatic cost optimization.
Set it once, save forever.

 Step 7: Query & Analytics**Query Engines:**

- Databricks SQL (serverless, auto-scaling)
- BigQuery (GCP serverless)
- Redshift (AWS-native)
- Athena (query S3 directly, cheapest)

BI Tools:

- Tableau (dashboards)
- Looker (embedded analytics)
- Power BI (Microsoft stack)
- Databricks SQL Analytics (built-in)

ML Access:

- Databricks Feature Store (native integration)
- Direct S3/GCS access (Spark, Python)

Query optimization:

- Delta Lake (Z-ordering, data skipping)
- Partition pruning (only scan needed dates)
- Column pruning (only read needed columns)
- Liquid clustering (automatic optimization)

Gold layer = optimized for this.

Queries run in **seconds**, not minutes.

Step 8: Data Observability

Track:

- Data freshness (is today's data here?)
- Volume anomalies (sudden spike/drop?)
- Schema changes (breaking downstream?)
- Null rates (data quality degrading?)
- Processing lag (how far behind?)

Tools:

- Great Expectations (data validation)
- Monte Carlo (data observability)
- dbt (data transformation + testing)
- Custom metrics (Prometheus + Grafana)

Alerts:

- Data missing (expected 10TB, got 1TB)
- Pipeline failed (Airflow DAG error)
- Query slow (Databricks warehouse overloaded)
- Cost spike (someone left Spark cluster running)

Without monitoring:

- Bad data propagates to Gold
- BI dashboards show wrong numbers
- ML models train on garbage
- Business makes decisions on bad data

Data quality > data quantity

⚠ Step 9: Design for Failure

Where does petabyte pipeline break?

Ingestion failures:

- ✗ → API rate limits (retry with backoff)
- ✗ → Source downtime (queue and replay)
- ✗ → Schema changes (alert, don't fail)

Processing failures:

- ✗ → OOM errors (tune Spark memory)
- ✗ → Skewed partitions (repartition)
- ✗ → Corrupt data (skip, log, alert)

Storage failures:

- ✗ → S3 throttling (batch smaller)
- ✗ → Disk full (auto-scale, lifecycle policies)

Cost explosions:

- ✗ → Runaway Spark jobs (set timeouts)
- ✗ → Full table scans (partition pruning)

Then design defenses:

- Idempotency (rerun safely)
- Checkpointing (resume where failed)
- Dead letter queues (bad data isolated)
- Circuit breakers (stop before it gets worse)

Top pipelines **fail gracefully**. And **recover automatically**.

Step 10: Cost Control at Scale

Cost optimization isn't optional at **petabyte scale**.
It's survival.

Where money goes:

Storage: 3.65 PB/year

→ Without tiering: **\$84K/year**

→ With tiering: **\$40K/year (50% savings)**

Compute (Databricks):

→ All-purpose clusters: **\$100K/year**

→ Jobs clusters (auto-terminate): **\$30K/year (70% savings)**

Queries (Databricks SQL):

→ Always-on warehouse: **\$50K/year**

→ Serverless SQL (auto-suspend): **\$20K/year (60% savings)**

Data transfer:

→ Cross-region: **\$10K/year**

→ Same region: **\$1K/year (90% savings)**

Total without optimization: \$244K/year

Total with optimization: \$91K/year

Savings: \$153K/year (63% reduction)